




OS TASK 5 - SHARED MEMORY & PTHREADS

Farah Kabesh - 900191706
Karim Sherif - 900191474





ROLES OF TEAM MEMBERS

Farah

- ❑ parallel_sum_mmap function.
- ❑ Excel File.
- ❑ Presentation.

Karim

- ❑ parallel_sum_threads function.
- ❑ Excel File.
- ❑ Presentation.



DESCRIPTION OF PARALLEL_SUM_MMAP

- ❑ `mmap()`: maps between the address space of the process to the memory object represented by the file descriptor.
- ❑ `mmap(NULL, sizeof(unsigned long long int), PROT_READ | PROT_WRITE, MAP_SHARED | MAP_ANONYMOUS, 0,0)`
- ❑ `PROT_READ & PROT_WRITE`: data can be read and written.



DESCRIPTION OF PARALLEL_SUM_MMAP

- ❑ MAP_SHARED & MAP_ANONYMOUS: shares mapping with all other processes and mapping is not connected to any files.
- ❑ File descriptor and offset value set to 0 because we have not mapped any file.
- ❑ Children write partial sums to shared memory.
- ❑ Parent waits for children to complete before reading the shared memory and computing the total sum.
- ❑ To unmap the mapped region (remove a mapping), use munmap().



PSEUDO CODE OF PARALLEL_SUM_MMAP PART 1

```
parallel_sum_mmap(unsigned int n_proc, unsigned int n)
```

Initialize:

```
s=0, partial=0
```

```
start, end=1
```

```
interval = n/n_proc
```

```
*shared = mmap(NULL, sizeof(unsigned long long int),  
PROT_READ | PROT_WRITE, MAP_SHARED | MAP_ANONYMOUS, 0,0)
```



PSEUDO CODE OF PARALLEL_SUM_MMAP

PART 2

```
if shared is equal to MAP_FAILED  
    print "mapping failed"  
    return 1
```

```
for i=1 to i <= n_proc  
    start = end  
    end = end + interval
```

```
if i is equal to n_proc  
    end = n + 1
```



PSEUDO CODE OF PARALLEL_SUM_MMAP

PART 3

```
if fork() is equal to 0
    for start to start < end
        partial = partial + start
        shared[i] = partial

    exit
else
    wait(NULL)
    s = s + shared[i]
```



PSEUDO CODE OF PARALLEL_SUM_MMAP PART 4

Initialize:

```
err = munmap(shared, sizeof(unsigned long long int))
```

```
if err not equal to 0
```

```
    print Unmapping Failed
```

```
    return 1
```

```
return s
```



DESCRIPTION OF PARALLEL_SUM_THREADS

- ❑ This program is a multithreaded process, where each thread calculates a partial sum and adds it to the total sum.
- ❑ `pthread_create` takes 4 parameters and creates a new thread.
- ❑ First argument is `pthread_t` which is used to store id of the thread
- ❑ Second argument is `attr` is pointer to struct containing the scheduling information.
- ❑ Third argument contains the function to be threaded.
- ❑ Fourth argument contains the argument for the the threaded function





DESCRIPTION OF PARALLEL_SUM_THREADS

- ❑ pthread_join is a function that lets process to wait for all the threads to finish.
- ❑ It takes the thread_id of the target thread and copies the exit status of that target.
- ❑ Mutexes is used to lock and unlock the final sum so no two threads can change its value at the same time.



PSEUDO CODE OF PARALLEL_SUM_THREADS

PART 1

Initialize:

sum_t

temp

struct thread_data

i, s, e, n, num_thread

pthread_mutex_t mutex1



PSEUDO CODE OF PARALLEL_SUM_THREADS

PART 2

```
void * threaded_sums(void* arg)
```

```
Initialize partial=0
```

```
if data->i is equal to data->num_thread  
    data->e = data->n
```

```
for data->s to data->s <= data->e  
    partial = partial + data->s
```



PSEUDO CODE OF PARALLEL_SUM_THREADS

PART 3

```
pthread_mutex_lock(&mutex1)
```

```
sum_t = sum_t + partial
```

```
if data->i is equal to data->num_thread  
    temp = sum_t
```

```
pthread_mutex_unlock(&mutex1)
```

```
return (void*) temp
```



PSEUDO CODE OF PARALLEL_SUM_THREADS

PART 4

```
parallel_sum_threads(n_thread, n)
```

Initialize:

```
start=1, end=interval, ret, *status
```

```
interval = n/n_thread
```

```
struct thread_data ta[n_thread]
```

```
pthread_t threads[n_thread]
```



PSEUDO CODE OF PARALLEL_SUM_THREADS PART 5

```
for j=0 to j < n_thread
```

```
    ta[j].i = j + 1
```

```
    ta[j].n = N
```

```
    ta[j].num_thread = n_thread
```

```
    ta[j].s = start
```

```
    ta[j].e = end
```

```
ret = pthread_create(&threads[j], NULL, threaded_sums,  
    (void *) &ta[j])
```



PSEUDO CODE OF PARALLEL_SUM_THREADS

PART 6

```
if ret not equal to 0  
    print "Thread can't be created"
```

```
start = end + 1  
end = end + interval
```

```
for j=0 to j < n_thread  
    pthread_join(threads[j], &status)
```

```
return status
```



HOW TEST 1 WAS DESIGNED & WHY?

FIXING N_PROC & N_THREAD

- Machine Name: Intel® Core™ i7-6700 CPU @ 3.40GHz × 8.
- Machine Specifications:
 - Number of cores = 4.
 - Number of threads = 8.
- Therefore n_proc = 3 and n_threads = 2.
- In order to obtain readings:
 - for (n = 1; n < 1000000; n+=5000)



HOW TEST 2 WAS DESIGNED & WHY?

FIXING N

- $n = 100000$.
- In order to obtain readings:
 - for ($j = 1; j < 150; j++$)

Note: j represents number of processes/threads.






QUESTION 1

If we fix n_{proc} and n_{thread} to the same value, then which of the 4 functions take the least time? At which value of N does this behavior persist?

GRAPH 1



ANSWER TO QUESTION 1

- ❑ Out of the 4 functions, the pipes function takes the least time.
- ❑ This behavior persists at the value $N = 75001$.
- ❑ Since n increases and the number of processes/threads is the same, time taken for execution using pipes and shared memory is less than time taken using sum and threads.

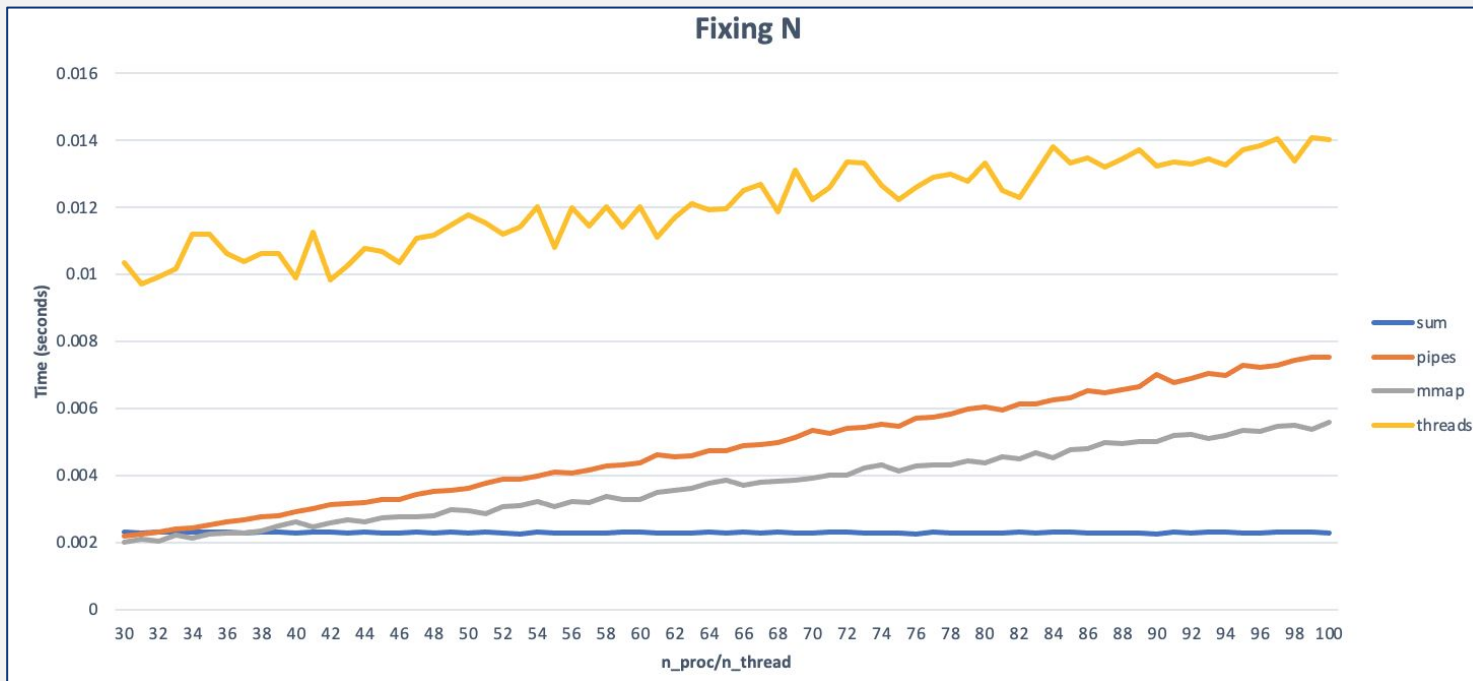




QUESTION 2

If we fix N to a large number, then which of the 4 functions take the least time? At which value of $n_{\text{proc}}/n_{\text{thread}}$ does this behavior persist?

GRAPH 2



ANSWER TO QUESTION 2

- ❑ Out of the 4 functions, the sum function takes the least time.
- ❑ This behavior persists at the value $n_proc/n_thread = 35$
- ❑ As n_proc increases, the overhead of creating new processes increases which leads to more time needed to compute the sums. Therefore, time taken for execution using sum is takes the least time; followed by memory, pipes, and finally threads.

